

Class 7: Machine Learning 1

Grace Wang (PID: A16968688)

Table of contents

Clustering	1
K-means	3
Hierarchical clustering	7
Dimensionality reduction - principal component analysis (PCA)	9
Using PCA	12
Interpreting PCA results	13


Today we will explore unsupervised machine learning methods, beginning with clustering and dimensionality reduction.

Clustering

Let's make up some data to cluster so that we know what the answer should be. The `rnorm()` function will help here.

```
hist(rnorm(10000, mean = 3))
```

Histogram of rnorm(10000, mean = 3)



Return 30 numbers centered on -3.

```
rnorm(30, mean = -3)
```

```
[1] -3.1334729 -3.2540254 -2.2029798 -0.4556317 -1.8410348 -3.3900962
[7] -3.7773526 -3.5013933 -2.8390137 -2.0760867 -4.7931490 -2.8556319
[13] -1.6098052 -3.6848046 -1.4995771 -2.9471420 -2.6221567 -2.2866287
[19] -0.1237786 -4.4097363 -4.1454898 -2.4952787 -2.6383429 -3.0258054
[25] -3.1326430 -3.5473095 -2.8748996 -3.3316792 -4.2395189 -1.9419915
```

```
rnorm(30, mean = 3)
```


```
[1] 2.580990 2.892710 3.407048 2.768087 3.026499 3.886969 3.049518 3.210014
[9] 2.720161 2.564673 3.412345 3.047199 2.540965 4.129591 4.066433 1.961232
[17] 1.426716 2.330056 3.526165 3.603371 2.380431 3.400538 2.121846 3.171463
[25] 3.667982 2.479904 3.204630 2.729300 3.483269 3.229509
```

```
tmp <- c(rnorm(30, mean = -3),
         rnorm(30, mean = 3))

x <- cbind(x = tmp, y = rev(tmp))
#x
```

Make a plot of `x`.

```
plot(x)
```



K-means

The main function in base R for K-means clustering is called `kmeans()`.

```
km <- kmeans(x, centers = 2)
```

The `kmeans` function returns a list with 9 components. You can see the named components of any list with the `attributes()` function.

```
attributes(km)
```

```
$names
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"

$class
[1] "kmeans"
```

How many points are in each cluster?

```
km$size
```

```
[1] 30 30
```

Cluster assignment/membership vector

```
km$cluster
```

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


Cluster centers

```
km$centers
```

```
      x      y
1  2.926942 -3.036133
2 -3.036133  2.926942
```

Make a plot of our `kmeans()` results showing cluster assignment using different colors for each group of points and cluster centers in blue. .


```
plot(x, col = km$cluster, pch = 19)
points(km$centers, col = "blue", pch = 15, cex = 2)
```



Run `kmeans()` again on `x` but cluster into 4 groups and plot the same result figure.


```
km2 <- kmeans(x, centers = 4)

plot(x, col = km2$cluster, pch = 19)
points(km2$centers, col = "blue", pch = 15, cex = 2)
```



Scree plot

```
within <- numeric(15)
for (n in 1:15){
  km <- kmeans(x, centers = n)
  within[n] <- km$tot.withinss
}
plot(x = 1:15, y = within)
lines(x = 1:15, y = within)
```



Key point - k-means clustering is very popular but can be misused. One big limitation is that it can impose a clustering pattern on data even if clear natural groupings don't exist - i.e., it does what you tell it to do with regard to the number of centers.

Hierarchical clustering

The main function in base R for hierarchical clustering is `hclust()`.

You can't just pass an input dataset as is into `hclust()`. Data must first be made into a "distance matrix", which can be done using the `dist()` function.

```
d <- dist(x)
hc <- hclust(d)
hc
```


Call:

```
hclust(d = d)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

The results of `hclust()` don't have a useful `print` method, but they do have a special `plot` method.

```
plot(hc)
abline(h = 8, col = "red")
```




To get our main cluster assignment (membership vector), we need to “cut” the tree at the big goalposts.

```
grps <- cutree(hc, h = 8)
table(grps)
```

```
grps
 1  2
30 30
```

```
plot(x, col = grps, pch = 19)
```

Hierarchical clustering is distinct in that the dendrogram can reveal the potential grouping in your data (unlike k-means).

Dimensionality reduction - principal component analysis (PCA)

PCA is a common and highly useful dimensionality reduction technique used in many fields - particularly bioinformatics.

Here we will analyze some data from the UK on food consumption.

```
UK_foods <- read.csv(url("https://tinyurl.com/UK-foods"))
dim(UK_foods)
```

```
[1] 17  5
```

```
head(UK_foods)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586

4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139


```
rownames(UK_foods) <- UK_foods[, 1]
UK_foods <- UK_foods[, -1]
head(UK_foods)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

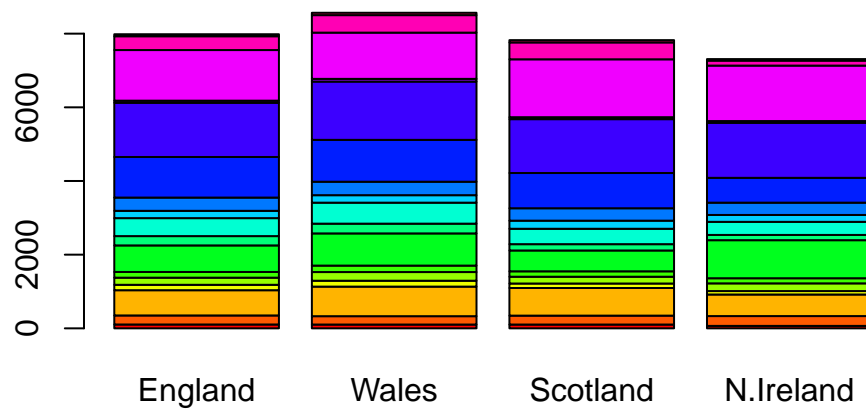
```
UK_foods <- read.csv(url("https://tinyurl.com/UK-foods"), row.names = 1)
head(UK_foods)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

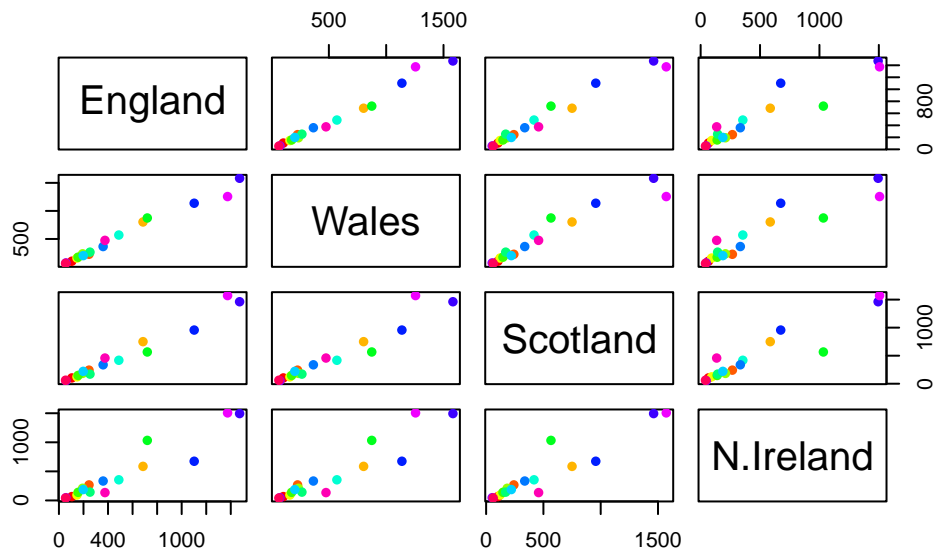
```
barplot(as.matrix(UK_foods), beside=T, col=rainbow(nrow(UK_foods)))
```



```
barplot(as.matrix(UK_foods), beside=F, col=rainbow(nrow(UK_foods)))
```



```
pairs(UK_foods, col = rainbow(nrow(UK_foods)), pch = 16)
```



Using PCA

The main function in base R for PCA is `prcomp()`.

```
UK_foods <- t(UK_foods)
pca <- prcomp(UK_foods)
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

The `prcomp()` function returns a list object of our results with 5 attributes.

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation" "center"    "scale"     "x"

$class
[1] "prcomp"
```

Interpreting PCA results

The two main “results” are `pca$x` and `pca$rotation`. The first of these contains the scores of the data on the new PC axis - we use these to make our PCA plot.

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-4.894696e-14
Wales	-240.52915	-224.646925	-56.475555	5.700024e-13
Scotland	-91.86934	286.081786	-44.415495	-7.460785e-13
N.Ireland	477.39164	-58.901862	-4.877895	2.321303e-13


```
library(ggplot2)
```

Warning: package 'ggplot2' was built under R version 4.4.3

```
library(ggrepel)
```

Warning: package 'ggrepel' was built under R version 4.4.3

```
# Make a plot of pca$x with PC1 vs PC2
ggplot(pca$x) +
  aes(x = PC1, y = PC2, label = rownames(pca$x)) +
  geom_point() +
  geom_text_repel()
```



The plot shows that England, Scotland, and Wales cluster together along the PC1 axis while Northern Ireland is on its own. England, Scotland and Wales are further apart along the PC2 axis, but the PC1 axis captures the largest proportion of the variance, so there is a larger difference between Northern Ireland and the rest of the countries.


The second major result is contained in the `pca$rotation` object or component. Let's plot this to see what PCA is picking up.

```
pca$rotation
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.694538519
Carcass_meat	0.047927628	0.013915823	0.06367111	0.489884628
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.279023718
Fish	-0.084414983	-0.050754947	0.03906481	-0.008483145
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257	0.076097502
Sugars	-0.037620983	-0.043021699	-0.03605745	0.034101334
Fresh_potatoes	0.401402060	-0.715017078	-0.20668248	-0.090972715
Fresh_Veg	-0.151849942	-0.144900268	0.21382237	-0.039901917
Other_Veg	-0.243593729	-0.225450923	-0.05332841	0.016719075
Processed_potatoes	-0.026886233	0.042850761	-0.07364902	0.030125166
Processed_Veg	-0.036488269	-0.045451802	0.05289191	-0.013969507
Fresh_fruit	-0.632640898	-0.177740743	0.40012865	0.184072217

Cereals	-0.047702858	-0.212599678	-0.35884921	0.191926714
Beverages	-0.026187756	-0.030560542	-0.04135860	0.004831876
Soft_drinks	0.232244140	0.555124311	-0.16942648	0.103508492
Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320	-0.316290619
Confectionery	-0.029650201	0.005949921	-0.05232164	0.001847469

```
ggplot(pca$rotation) +
  aes(x = PC1, rownames(pca$rotation)) +
  geom_col()
```



The figure shows how each variable contributes to PC1. We can see that positive values for PC1 are associated with more soft drink and fresh potato consumption and less fresh fruit and alcohol consumption.